ATLASSIAN

Atlassian's Submission to the Australian Government in relation to Safe and Responsible AI in Australia

Technology Strategy Branch
Department of Industry, Science and Resources
DigitalEconomy@industry.gov.au

4 August 2023

We appreciate this opportunity to provide feedback to the Department on appropriate governance, regulatory and policy mechanisms to facilitate the safe and responsible development and use of AI in Australia, as set out in the Discussion Paper of June 2023 (the **Discussion Paper**).

At Atlassian, we build enterprise software products to help teams collaborate, including for software development, project management and content management. As one of Australia's most successful home-grown technology companies — and one that provides products and services to customers around the world — we believe that we are in a unique position to contribute to this consultation.

We know the critical role that technology (including emerging technology) plays in powering the operations of our customers, and the digital economy more broadly. Our customers' technology engineering teams around the world use our products as key tools that help drive their development of technology products and services throughout their lifecycle.

Like many companies, we also leverage the power of machine learning in our own products. For years we have used machine learning to enhance core experiences in our products, from personalised search to recommending people and teams to bring into collaboration. This April, we also announced Atlassian Intelligence, a step forward in bringing the power of AI (including large language models) to our full family of cloud products.[1] We also affirmed our commitment to working with and helping our customers and partners navigate this fast-changing technology landscape responsibly and in line with the values that we all share, in line with our Responsible Technology Principles.[2]

However, as a company that specialises in collaboration and teamwork, we know that individual company principles and commitments aren't enough. We all have a role to play in fostering a thriving, trusted emerging technology ecosystem.

We believe that our regulatory landscape needs to clearly and carefully anticipate what our digital future will mean for Australian organisations and individuals operating in an evolving global economy. This consultation process presents an ideal opportunity to carefully assess the current state of our digital regulatory framework, and to set forth a practical, forward-looking roadmap for responding to the opportunities and issues raised by emerging technologies like AI in a clear, considered and holistic manner.

[1] See https://www.atlassian.com/software/artificial-intelligence.

[2] See https://www.atlassian.com/trust/responsible-tech-principles.

Our approach and proposed model

In late 2020, Atlassian published eight Principles for Sound Tech Policy.[3] These Principles are intended to not only guide Atlassian's own engagement on important matters of public policy, but to set forth guiding principles for what we believe sound technology-related public policy should look like more broadly.

Atlassian has also been proud to contribute to the Tech Policy Design Kit developed by the ANU Tech Policy Design Centre, in collaboration with the Tech Council of Australia (of which Atlassian is a founding member) and the Digital Technology Taskforce.[4] The Tech Policy Design Kit is intended to provide a foundation for designing effective technology policy.

In line with these Principles and processes, we strongly recommend that the Australian Government use this opportunity to design and set forth a purpose-driven, outcomes-focused approach to the governance of AI and other emerging technologies.

Ultimately, we believe that the best response to these emerging technologies is situated within an overarching, coordinated digital regulatory framework that is:

- governed by **core principles**, which may be enshrined in legislation and would set forth a consistent, scalable and risk-based framework for all stakeholders, and inform the formulation and implementation of specific measures and tools within that framework;
- supported through one or more central **advisory bodies** that are capable of providing advice and assisting with coordination and alignment across government agencies and regulators with responsibility across various sectors and areas of law, which would also allow government to build expertise (as to how technology operates, the opportunities and challenges it creates and how best to respond) and connections with industry; and
- bolstered by an appropriate mix of targeted and objective **governance measures, guidance and tools** (including regulatory measures where required), which respond clearly to identified issues in a manner that aligns to and has the benefit of the overarching principles and institutional expertise.

In our view, this proposed model has the flexibility to account for the multi-dimensional nature of many of the opportunities, risks and issues that can be raised by AI. It also allows the governance model to be imbued with Australian values and priorities in its overarching principles, while also retaining the flexibility to both interoperate with and take advantage of emerging international standards and best practice through its agile approach to different policy and regulatory responses.

This recommended overarching model is, of course, only a starting point. In this submission, we set forth our vision not only for *what* framework should be implemented, but also *how* it should be achieved (and *why*).

Our recommended roadmap: From framework to action

1. Define the playing field

The public release of OpenAI's ChatGPT in late 2022 — and the incredible speed of its adoption by a vast and varied user base — has helped to demonstrate the capabilities and the potential of AI to an audience that may not have appreciated just how far AI technologies have already advanced.

But this perception of swift and unprecedented advancement means that just as many individuals and consumers are now discovering the promise and opportunities of AI, they are

[3] These Principles are also available for download at https://www.atlassian.com/blog/technology/regulating-technology.

[4] See https://techpolicydesign.au/tech-policy-design-kit.

also discovering the actual and potential issues and harms associated with its use and adoption.

These issues and harms are, in many respects, not new. Instead, the scale and speed involved has led to a perception that we are now treading in territory that is as new and uncharted from a legal and regulatory perspective as it is from a technological one.

This is not the case. But we believe that it does illustrate an urgent and critical need to build a common understanding of how the law does — and should — apply to AI and other emerging technologies.

First, we acknowledge that the issues and harms that may be raised by AI can be complex. As the Discussion Paper makes clear, AI is not a singular technology but an umbrella term for a constellation of technologies, which can involve the use and application of a variety of techniques, processes and tools across a broad spectrum of sectors, contexts and use cases. These range from the productivity solutions that individuals and businesses use on a daily basis (including Atlassian's products) to highly specialised and complex use cases within a range of sectors.

The breadth of these use cases and their application may then give rise to varying, context-dependent risk profiles with legal, regulatory and broader societal dimensions.

Further, we believe that law and regulation, designed carefully, can be a 'force multiplier' for trust: fostering confidence and trust in new industries and technologies, and encouraging their adoption in a way that brings the most benefit to us all.

However, in many contexts, the use of AI is already subject to existing laws and frameworks that apply equally to AI or might be capable of being extended to AI systems under certain conditions. The Discussion Paper outlines some of these examples, and the Human Technology Institute (**HTI**) has set forth a more detailed overview of the landscape in its May 2023 report on the State of AI Governance in Australia.[5]

What is not always clear from these outlines and overviews is *why* 'the law' as it stands is not currently being applied, or cannot be applied, to AI. This is because there are a number of different factors that may be at play including, among others:

- where laws do apply, but are not being enforced (for example, due to a lack of awareness or resources) or are insufficient on their own to drive the right outcomes;
- where laws may apply, subject to further clarifications and guidance; and
- where laws could or should apply, with the application of reforms — which may:
    - already be contemplated as part of processes that are underway (such as the Privacy Act review); and
    - span a spectrum from minor reforms based on first principles (for example, to acknowledge the presence of automated decision-making processes in place of human decision-makers) to major reforms needed to account for actual harms that are not currently addressed by our legal and regulatory frameworks.

Understanding this context is a critical first step to establishing a clear governance model for AI in Australia, as it provides the foundations for understanding the role of law and regulation in responding to emerging technologies like AI — and for implementing a framework that is responsive to the associated issues.

[5] See https://www.uts.edu.au/human-technology-institute/projects/ai-corporate-governance. Atlassian is a proud advisory partner of HTI and is excited to continue to partner with HTI on its new program of work, The Future of AI Regulation in Australia, which we believe will also be of significant benefit to the goals of this consultation.

2. Engage with the issue

From the baseline of this common understanding of where we are today and informed by the core principles at the centre of our proposed model, we can then consider how to get to where we want to be.

In particular, the above breakdown of our current landscape can form the basis for prioritising the tools and measures needed to drive towards that outcome. At this relatively early stage, that prioritisation may look like a broad, 'traffic-light' overview of the landscape:

(1) **Red:** Scenarios of actual, present or imminent harms, for which appropriate interventions don't currently exist and where a proactive approach is justified.

(2) **Amber:** Scenarios where it is not clear whether appropriate interventions already exist, or whether these may need to be bolstered by further reactive measures. In this case, more long-term work may be needed to identify and implement a program of work and (where relevant) reforms.

(3) **Green:** Scenarios where laws currently apply or would apply without the need for reform or investigation, but where non-regulatory tools (including greater awareness, resources and guidance) would assist to build and maintain public trust in safe and responsible AI.

We appreciate that it may be difficult to definitively identify which category each use case or scenario may fall into at this stage. This categorisation is accordingly intended to be fluid rather than fixed, and would benefit significantly from the institutional expertise of the advisory body (or bodies) forming part of our proposed model as relevant issues, harms and opportunities emerge and change over time.

3. Treat the ailment, don't kill the patient

Based on the above categorisation and prioritisation, regulatory measures that are targeted towards prevention and remediation of identified harms can then be designed and implemented where needed.

Initially focused on the red category, with a view towards addressing amber items at later stages, these measures should flow from the principles-based framework and be:

- risk-based by their nature, having evolved from the context of use and application, as well as the categorisation and prioritisation of actual and potential harms involved;
- targeted at addressing where and how the problem or risk under consideration arises, its relationship to the relevant AI technologies and those who provide, implement or use them (as applicable);
- technology-neutral, unless there is a clear and valid justification for implementing a technology-specific measure; and
- with the benefit of the advisory expertise provided above, cognisant of emerging international practice and standards in the circumstances, with a bias towards being interoperable with those standards.

As noted above, while several red items may be capable of being identified and responded to in earlier stages, we appreciate that amber items are likely to require greater effort to identify and assess, and may still emerge over time as existing concepts and norms are tested by emerging technologies.

Important work has already been done and will continue to be done in this space, including by HTI. However, the more complex and evolving nature of these issues is such that they would benefit from a more comprehensive review of our current legal framework to ensure it addresses the implications of widespread adoption of AI and emerging technologies.

This longer-term project could be undertaken by one or more of the advisory bodies under our proposed model, or by an existing body with the expertise and capacity to undertake

these sorts of reviews like the Australian Law Reform Commission. While by its nature this review would be expected to take time, the terms of reference for the review could account for interim reports and the capacity to consult and advise on urgent legal issues and questions as they arise — as AI technologies are diffused across the Australian economy and society — striking the balance between delivering longer-term certainty and remaining flexible and responsive to emerging concerns.

4. Tech (and trust) is global

Importantly, and as noted in the Discussion Paper, available regulatory and policy options do not have to be limited only to 'hard' law and regulation. The most effective approach is likely to be composed of a broad and flexible mix of targeted 'hard' regulatory measures outlined above, as well as 'soft' law, standards, resourcing and investment, to collectively build a culture of safe and responsible AI in Australia.

In this respect, emerging international responses are illustrative. In particular, many overseas models are demonstrating an increased reliance upon technical and industry standards to drive consistency and uplift industry practice. This includes the US National Institute of Science and Technology's AI Risk Management Framework, the UK AI Standards Hub and references to standards and specifications throughout the EU AI Act (all referenced in the Discussion Paper).

These examples are not only useful as learnings for Australia in which measures to consider implementing locally. The inherently global nature of technology (including AI) is such that emerging global standards and best practice can and will have a broader normative effect.

This means that there is a clear opportunity for Australia to:

- invest more heavily in standards development processes — ensuring that its voice and values are accounted for in these standards; and
- drive alignment and interoperability with these standards — helping to build a trusted ecosystem of compliance and certification (on the basis that adopting common or harmonised standards would open up access for both providers and consumers to a wider range of certified providers and certification bodies globally).

5. Build the foundation for shared success

This mix of approaches will not only be critical to addressing the green category above, but also essential to the broader success of the overall model.

In order for Australians (and Australia) to fully and equitably harness the benefits of AI — and be capable of exercising their rights in the digital economy of tomorrow — we will therefore need significant investment in education, skills and awareness.

Case study: To illustrate why this is the case, consider this example:

- In some cases, individuals are not made aware that they are interacting with 'human-like' AI systems (like chatbots), or that AI technologies are being used to make certain significant decisions about their lives and livelihoods (like eligibility for insurance).
- Given the potential harms, these scenarios are likely to fall within the red category, and new mandatory notification requirements would need to be designed in response.
- However, if individuals receiving that notification aren't equipped with an understanding of what it means when an AI system is in use, how best to interact with it and how to respond to (and if need be, challenge) its outputs, then those notification requirements could be effectively meaningless.

This investment requires a broad view of lifelong education efforts, with a focus on critical thinking and interdisciplinary skills, aiming to encourage curiosity and discourage tendencies towards unquestioning acceptance and automation bias.

It therefore needs to include education and awareness-raising at the levels of:

- regulators, policymakers and government agencies, with the benefit of the institutional expertise of the advisory body (or bodies) forming part of our proposed model, in order to ensure that our laws can be applied and enforced consistently in the age of AI;
- within the workforce, to encourage responsible governance of AI throughout all levels of organisations (including in line with the recommendations set forth in HTI's report on the State of AI Governance in Australia);
- in our education and training system, so that we raise the awareness of AI use cases, that these skills are embedded as early as possible and that young Australians are well-equipped for our digital future.

While this is no small task, we believe that now is the ideal time to set forth a comprehensive and multi-dimensional vision and long-term roadmap for realising the opportunities and minimising the harms of AI in Australia.

Atlassian would be pleased to further discuss these comments with the Department, and looks forward to working with the Australian Government to ensure that all Australians will benefit from our AI-powered future.

Yours sincerely,

Anna Jaffe
Director of Regulatory Affairs & Ethics
Atlassian

David Masters
Head of Global Policy & Regulatory Affairs
Atlassian