

18 August 2023

Executive summary

Canva is an Australian founded and headquartered global online visual communication platform that serves 135 million people all over the world. Founded on the belief that people shouldn't need to understand complex software to unlock their creativity, we've set ourselves a mission to empower everyone to design, accelerating this through supercharging our Visual Suite with Artificial Intelligence (AI)-powered capabilities.¹

Through this submission, we aim to provide insight into critical areas, advocating for policies that will promote safe, ethical, and responsible use of AI. Broadly, we believe in:

- **International alignment on definitions and standards for AI:** By fostering agreement on key definitions and standards, we can ensure consistency and mutual understanding across borders, making cooperation and collaboration more feasible.
- **Interoperability across AI governance frameworks:** Given the global nature of AI technology, having unified policies and regulations will ensure the safe use of AI and simplify cross-border operations to promote tech growth.
- **Transparency in AI systems:** Ensuring that AI systems operate transparently is key to building trust and mitigating potential risks or misunderstandings. We also believe in the importance of helping the public interact with AI in an informed way. This involves consumer protections, such as watermarking, which can ensure authenticity and provenance of AI-generated content.

Canva's approach to AI

Canva's mission is to empower the world to design. We prioritise responsibility, simplicity, and creativity in our products, community interactions, and AI development.

We view AI as a collaborative partner in the creative process. As we strive to make design accessible to all, AI can play a pivotal role in enabling individuals to transform their ideas into reality without the need for specialised equipment or tools.¹ The core of the design process lies in human intention and creativity. While AI cannot independently create, its potential lies in amplifying the creativity and intentions of the human designer and helping people discover innovative starting points.

Section 1

Section 1. The way we use AI

Our Canva community is made up of solo entrepreneurs, small businesses, nonprofits, teachers, students, global enterprises and many more who've created over 15 billion designs on Canva since 2013. As content creation continues to grow, we believe it's our responsibility to continually evolve our policies and tools to support our community and the content they create.

At Canva, we utilise AI in the products and tools we provide. Some of our primary products include Text to Image², Magic Write™³, Magic Edit⁴, Magic Design⁵, and Translate.⁶

For example, we use Generative AI (GenAI) to create and transform content for use in designs.⁷

1. "Canva Create," Canva, accessed 31 July 2023, <https://www.canva.com/canva-create/>.

2. "AI Image Generator," Canva, accessed 3 August 2023, <https://www.canva.com/ai-image-generator/>.

3. "Magic Write," Canva, accessed 3 August 2023, <https://www.canva.com/magic-write/>.

4. "AI Photo Editing," Canva, accessed 3 August 2023, <https://www.canva.com/features/ai-photo-editing/>.

5. "Magic Design," Canva, accessed 31 July 2023, <https://www.canva.com/magic-design/>.

6. "Translate," Canva, accessed 3 August 2023, <https://www.canva.com/features/translate/>.

7. "AI Product Terms," Canva, accessed 3 August 2023, <https://www.canva.com/policies/ai-product-terms/>.

GenAI models create new content in response to user inputs based on abstract principles drawn from the data used to train the model.⁸

Non-Generative AI (non-GenAI) is used to optimise workflows and existing features by predicting recommended actions and streamlining user interactions. An example is our one-click Background Remover feature that has been used over two billion times. It removes backgrounds in seconds, rather than having to painstakingly spend time doing this manually.⁹

Furthermore, non-GenAI plays a crucial role in tailoring user experiences, safeguarding against fraud and other abuse, identifying bugs, and performing other essential functions. Both types of AI help us provide the best experience for all Canva users.

AI safety at Canva

Our values play an incredibly important role in guiding how we build our products, treat our community, and the impact we want to have on the world. We're committed to building a safe and trusted environment where our community can unlock their creativity and bring their ideas to life. In order to maintain a safe and respectful environment, we have implemented several trust and safety measures in our AI systems.¹⁰

First, we conduct automatic moderation on user inputs into our GenAI systems (including prompts) to proactively identify and limit the generation of inappropriate content. Additionally, we extend this moderation to the output generated by our GenAI, intervening when we detect potentially undesirable, offensive, or harmful results. Another critical aspect of our commitment is the mitigation of bias in our AI systems and datasets. We recognise that biases can influence the outcomes, and what may be considered fair can vary across cultures. To address this, we take deliberate steps to minimise unfair biases and avoid misrepresentation, so that our AI systems serve all users equitably.

8. Karine Perset, Audrey Plonk, and Stuart Russell, "As Language Models and Generative AI Take the World by Storm, the OECD Is Tracking the Policy Implications," OECD AI Policy Observatory, 13 April 2023, <https://oecd.ai/en/wonk/language-models-policy-implications>.

9. Canva Newsroom, "Canva Developers Platform," accessed 31 July 2023 <https://www.canva.com/newsroom/news/canva-developers-platform/>.

10. "AI Safety," Canva, accessed 3 August 2023, <https://www.canva.com/policies/ai-safety/>.

We also use a range of tools to help to keep our platform safe, including AI systems dedicated to the task.

For example, our tools help our community adhere to our AI Product Terms by aiming to limit the creation of content that includes medical topics, self harm and mental health, political topics, or sexually explicit content.¹¹

To maintain the highest level of safety and quality, our features undergo a rigorous safety review process, scrutinised to adhere to our strict standards and guidelines.

In addition to our normal processes, we employ safety testing of AI features and have a global Trust & Safety operations centre.

However, in instances where anything inappropriate or unexpected is generated, or if a prompt has been incorrectly flagged, we have the option for people to submit a report to us so we can look into the issue.

When using Text to Image, people can report undesirable content using the “Report” button in the app, or if they’re using any of Canva’s other products, they can use our dedicated reporting form in our Help Centre.¹²

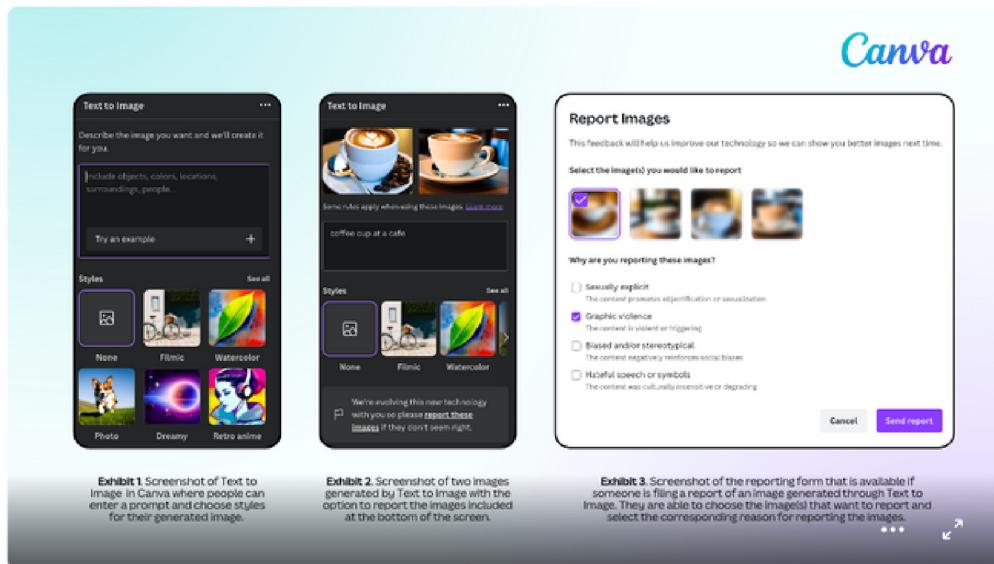


Figure 1. Example of reporting flow for Canva's Text to Image tool.

11. “AI Product Terms,” Canva, accessed 31 July 2023, <https://www.canva.com/policies/ai-product-terms/>.

12. Canva, “Report Content,” Canva Help Center, access 31 July 2023, https://www.canva.com/en_au/help/report-content/.

AI inclusivity at Canva

Inclusivity lies at the heart of the design of AI-powered capabilities, as we consciously design them to represent the diversity of our global user base. Informed by a team of in-house researchers who work to understand how our users around the world interact with AI, we encompass a broad spectrum of perspectives.

As mentioned, we are committed to the reduction of bias in our AI systems and datasets, as we recognise that biases can influence the outcomes, and what may be considered fair can vary across cultures. We believe in taking a responsive approach and we continuously develop new techniques for steering the bias in GenAI towards a fairer, more inclusive result.

AI fairness and inclusivity is an international concern, and Canva is engaged with both our industry peers and academia to share our findings and evolve both our technical and policy understanding of this space.

The convergence of advancing technology and the increasing demand for visual communication is unveiling previously untapped design capabilities that are already empowering people worldwide to unleash their creativity whilst unlocking productivity. That is precisely why we are introducing innovative AI-powered tools, and want to work collaboratively to harness the ways the world will use these innovations safely and responsibly.

Section 2



Section 2. Responses to discussion paper

In this consultation, we've organised our submission to respond to the questions in the discussion paper as they relate to Canva. We've provided insights in areas where we think we can make the most significant contribution, and we will continue to share our thoughts and work as the AI field keeps advancing and evolving.

Question 1: Do you agree with the definitions in this discussion paper? If not, what definitions do you prefer and why?

We agree that establishing uniform, global definitions for AI is crucial. Additionally, differentiating between non-GenAI and GenAI is essential given their different risks. This is important because:

- **We believe it's important to adopt global definitions** for technologies like AI, Machine Learning (ML), and algorithms, like those used by the Organisation for Economic Co-operation and Development (OECD), as this will help achieve a consistent approach across different nations and sectors.¹³ Standardised definitions for AI that are accepted worldwide can also support the global growth of Australian businesses by reducing barriers to entry and the need for significant operational adjustments or overhauls to meet external regulations in other countries.
- **It is also important to distinguish between non-GenAI and GenAI** due to their different risk factors.¹⁴ For example, non-GenAI tends to operate with a narrow and predefined set of instructions, and its behaviour is more predictable, like in the example of using AI to remove the background from an image. Conversely, GenAI is fundamentally based on random generation of new content based on human direction, which can be unpredictable, such as text generation in a model like GPT-4. Due to its specificity, non-GenAI's operations tend to be more transparent, whereas GenAI's underlying unpredictability and complexity means that understanding how and why it makes individual decisions is much more opaque. By clearly defining non-GenAI and GenAI, policymakers can craft more precise regulations that address the unique challenges and risks posed by each. This helps AI creators and users know their ethical and legal limits, encouraging safe AI practices.

Question 2: What potential risks from AI are not covered by Australia's existing regulatory approaches? Do you have suggestions for possible regulatory action to mitigate these risks?

13. OECD, "Artificial Intelligence & Responsible Business Conduct," OECD Guidelines for Multinational Enterprises, <https://mneguidelines.oecd.org/RBC-and-artificial-intelligence.pdf>.

14. McKinsey, "What is Generative AI?," McKinsey Explainers, <https://www.mckinsey.com/featured-insights/mckinsey-explainers/what-is-generative-ai>.

A significant challenge for businesses dealing with AI regulations is inconsistencies across different regions. Aligning with jurisdictions that share our views, such as the EU, would be a practical solution to address this issue. This is important because:

- **At a company like Canva that operates across 190 countries**, it's essential that AI policies are harmonised and interoperable to allow for smooth functioning across different countries.
- **AI systems are often made in one country and used in others.** Having consistent policies helps these systems work smoothly across borders and prevents "regulatory fragmentation" where different regions have very different rules. This fragmentation can make international AI development and use challenging for companies.
- **One way to help achieve this would be by considering a "safe harbour" system** with countries that have comparative or mutual AI safety frameworks.¹⁵ For instance, if Australian companies follow local rules and standards, they could also meet the EU's AI Act requirements when set. This way, Australian and EU AI startups can cooperate more easily, promoting better international AI cooperation and efficient policy application.
- **The same AI tool may be used across multiple organisations in multiple countries so it's important that they operate across borders.** This international reach means that the AI tools already receive regulatory scrutiny via different impact assessment methodologies, including the Data Protection Impact Assessments (DPIA) mandated by the EU's General Data Protection Regulation (GDPR)¹⁶ and various Algorithmic Impact Assessment approaches.¹⁷ Given the complexity of multiple legal requirements, we propose a tiered or graduated regulatory framework that will encourage growth with startups and Small and Medium-sized Businesses (SMBs), while also promoting ethical and responsible use of AI.

15. European Commission, "Press Corner - European Commission Statement," European Commission, accessed 3 August 2023, https://ec.europa.eu/commission/presscorner/detail/en/ip_23_3721.

16. GDPR.eu, "Data Protection Impact Assessment Template," GDPR.eu, [https://gdpr.eu/data-protection-impact-assessment-template/#:~:text=A%20Data%20Protection%20Impact%20Assessment%20\(DPIA\)%20is%20required%20under%20the,help%20you%20execute%20the%20assessment](https://gdpr.eu/data-protection-impact-assessment-template/#:~:text=A%20Data%20Protection%20Impact%20Assessment%20(DPIA)%20is%20required%20under%20the,help%20you%20execute%20the%20assessment).

17. OECD AI Policy Observatory, "Algorithmic Impact Assessment Tool," OECD, <https://oecd.ai/en/catalogue/tools/algorithmic-impact-assessment-tool>.

Question 3: Are there any further non-regulatory initiatives the Australian Government could implement to support responsible AI practices in Australia? Please describe these and their benefits or impacts.

There are a number of initiatives we'd like to see the Australian Government do to boost responsible AI practices, including allocating federal funding towards Research & Development (R&D) in AI technologies, with an emphasis on educational institutions and schools. This means:

- **Enable targeted government research grants.** Australia has the capability and desire to become a world leader in responsible AI, and targeted government research grants in this space would cement this outcome. This would offer incentives for top talent to conduct their AI-related research in Australia. More crucially, it would give local talent a compelling reason to remain in the country, rather than seeking opportunities in global hubs like Silicon Valley.

Question 4: Do you have suggestions on coordination on AI governance across government? Please outline any goals that any coordination mechanisms could achieve and how they could influence the development and uptake of AI in Australia.

The key change that the government could implement is cross-portfolio consistency across agencies and businesses. This means:

- **At present, Australia has multiple regulators with legislative functions related to AI.** As AI technology progresses, these overlaps can lead to complexity for businesses, including confusion and legal uncertainty.
- **Cross-portfolio consistency can provide both business and consumer certainty.** If government departments and agencies can act consistently, it will provide businesses with clarity on their compliance obligations.
- **Additionally, cross-agency coordination can clarify pathways for the general public to engage** or report AI-related concerns to the government, fostering public trust and enabling swift, coordinated responses to potential challenges.

- A cross-agency AI governance model should be formed with representatives from key agencies, such as the Department of Industry, Science, and Resources (DISR), the Office of the eSafety Commissioner (eSafety), the Australian Communications and Media Authority (ACMA), the Australian Competition and Consumer Commission (ACCC), the Australian Human Rights Commission (AHRC), and the Office of the Australian Information Commissioner (OAIC). This will ensure there are defined roles for each agency and prevent overlap.
- Additionally, any cross-agency governance model should include industry representatives to help establish channels of communication and coordination.

Question 6: Should different approaches apply to public and private sector use of AI technologies? If so, how should the approaches differ?

We believe there should be a focus on the purpose and risks of AI, rather than the difference in public or private use:

- Both the public and private sectors can use AI for various tasks, but the important thing to focus on is that **risks change based on how it's used**.
- **The agency and human rights of individuals affected by AI should ideally be the driving factor in how policy is applied.**
There are situations in which AI is used opaquely that span both the public and private sector, such as in physical security contexts (e.g. monitoring camera feeds for threats) and much more hidden contexts (e.g. making financial determinations about individuals).
- **Trust and individual agency is key in these opaque scenarios, and it is important to default to disclosing** that this kind of decision-making is occurring and being transparent about how it decides, and if possible offering avenues to opt-out or appeal these decisions. It's worth noting that such disclosures are crucial in public settings where sensitive data is involved, and individuals have fewer choices when using essential government services compared to the private sector, which has multiple AI options and alternatives.

Question 8: In what circumstances are generic solutions to the risks of AI most valuable? And in what circumstances are technology-specific solutions better? Please provide some examples.

AI is a broad class of technology, and the specific technologies and risks can be very different. We believe that a combination of general and technology-specific solutions are the most appropriate:

- Targeted solutions for specific types of AI, including both GenAI and non-GenAI, are better suited to address specific challenges due to the diverse nature of AI models, such as particular model weaknesses.
- General solutions to AI risks, like metadata and reporting mandates, can help with some of the basic ethical and legal issues that are common to many AI technologies, like setting transparency and provenance standards. However, specific solutions that are tailored to a certain types of AI are better for addressing unique risks, such as certain model vulnerabilities, and dataset-specific issues such as bias and fairness. General solutions provide a base layer of confidence in safety, but they must be coupled with specific solutions so we can use AI safely and responsibly as it keeps changing and growing.

Question 9: Given the importance of transparency in the AI lifecycle, please share your thoughts on: (1) where and when transparency will be most critical and valuable to mitigate potential AI risks and to improve public trust and confidence in AI; (2) mandating transparency requirements across the private and public sectors, including how these requirements could be implemented.

Transparency in the development and application of AI development is crucial for the ethical implementation of AI systems while also balancing the need for businesses to compete globally. We believe:

- Global collaboration and co-development of standards for signifying the authenticity of content, and whether it is fully AI-generated or partial, are imperative and provide transparency and data clarity.

- **Users should be informed about how their data is utilised**, such as whether it's used to train AI models, and be granted control options like opting in or opting out.
- **Transparency requirements should be tailored to the AI application's impact.** While AI decisions with significant consequences for individuals or businesses must be transparent, when AI is used for routine, low-risk functions like spam filtering, it may not require the same level of disclosure.

Question 11: What initiatives or government action can increase public trust in AI deployment to encourage more people to use AI?

To increase public trust, the Australian Government should promote AI literacy among people of all ages to increase awareness and understanding about both the capabilities and limitations of different AI systems. This could include:

- **Institute educational programs and workforce training on AI.** The Government can help people adjust to the growing influence of AI by introducing educational programs and workforce training on AI. Investing in AI education across all levels, from schools to universities and job training programs, can prepare the workforce to confidently use AI technologies.
- **Develop skills training.** Programs aimed at improving skills for AI-related jobs can also help workers stay competitive. Government and industry partnerships can create specific training for AI needs in different fields. Making training accessible and affordable, especially for underprivileged communities, is vital to ensure everyone can gain from AI education and opportunities. Encouraging continuous learning and adapting to AI's power can be achieved through incentives, grants, and collaborations with businesses. These steps can prepare a workforce that can innovate and flourish with AI technologies.

Question 14: Do you support a risk-based approach for addressing potential AI risks? If not, is there a better approach?

Canva supports an approach that looks at the different risk levels of different AI approaches and treats each risk appropriately. This is because:

- **AI and ML technology is already used in many ways, and the risk associated with an AI product varies widely depending on the application.** For instance, an AI tool designed solely to remove image backgrounds poses a much lower risk compared to a versatile model like GPT-4, which could potentially generate harmful content from benign prompts.
- **Open-source models are fundamental for research, development, innovation, competition, oversight, and safety improvements.** Underlying open-source models, which is a special category, should be subject to the same standards for rigour and safety. However, open-source models are not the same as commercial products or closed-source AI systems.¹⁸ To maintain global consistency, aligning with international standards like those set by the OECD, International Organization for Standardization (ISO), and the Hiroshima Process is crucial.¹⁹ These standards guide the responsible and ethical use of AI technologies worldwide. Modifying privacy and intellectual property laws can also balance fostering AI advancement with safeguarding individual and business rights.

18. GitHub Blog, "Supporting Open Source and Open Science in the EU AI Act," GitHub, accessed 3 August 2023, <https://github.blog/wp-content/uploads/2023/07/Supporting-Open-Source-and-Open-Science-in-the-EU-AI-Act.pdf>.

19. G7 Leaders. "G7 Hiroshima Leaders' Communiqué." The White House, May 20, 2023, <https://www.whitehouse.gov/briefing-room/statements-releases/2023/05/20/g7-hiroshima-leaders-communique/>.